



# A trip to Asymptopia

Statistical Inference

Brian Caffo, Jeff Leek, Roger Peng  
Johns Hopkins Bloomberg School of Public Health

```
## Error: object 'opts_chunk' not found
```

```
## Error: object 'knit_hooks' not found
```

```
## Error: object 'knit_hooks' not found
```

# Asymptotics

- Asymptotics is the term for the behavior of statistics as the sample size (or some other relevant quantity) limits to infinity (or some other relevant number)
- (Asymptopia is my name for the land of asymptotics, where everything works out well and there's no messes. The land of infinite data is nice that way.)
- Asymptotics are incredibly useful for simple statistical inference and approximations
- (Not covered in this class) Asymptotics often lead to nice understanding of procedures
- Asymptotics generally give no assurances about finite sample performance
  - The kinds of asymptotics that do are orders of magnitude more difficult to work with

Asymptotics form the basis for frequency interpretation of probabilities (the long run proportion of

# Numerical limits

- Imagine a sequence
  - $a_1 = .9$ ,
  - $a_2 = .99$ ,
  - $a_3 = .999$ , ...
- Clearly this sequence converges to 1
- Definition of a limit: For any fixed distance we can find a point in the sequence so that the sequence is closer to the limit than that distance from that point on

# Limits of random variables


- The problem is harder for random variables
- Consider  $\bar{X}_n$  the sample average of the first  $n$  of a collection of iid observations
  - Example  $\bar{X}_n$  could be the average of the result of  $n$  coin flips (i.e. the sample proportion of heads)
- We say that  $\bar{X}_n$  converges in probability to a limit if for any fixed distance the probability of  $\bar{X}_n$  being closer (further away) than that distance from the limit converges to one (zero)

# The Law of Large Numbers

- Establishing that a random sequence converges to a limit is hard
- Fortunately, we have a theorem that does all the work for us, called the **Law of Large Numbers**
- The law of large numbers states that if  $X_1, \dots, X_n$  are iid from a population with mean  $\mu$  and variance  $\sigma^2$  then  $\bar{X}_n$  converges in probability to  $\mu$
- (There are many variations on the LLN; we are using a particularly lazy version, my favorite kind of version)

# Law of large numbers in action

```
n <- 10000  
means <- cumsum(rnorm(n))/(1:n)  
plot(1:n, means, type = "l", lwd = 2, frame = FALSE, ylab = "cumulative means", xlab = "sample size")  
abline(h = 0)
```



# Discussion

- An estimator is **consistent** if it converges to what you want to estimate
  - Consistency is neither necessary nor sufficient for one estimator to be better than another
  - Typically, good estimators are consistent; it's not too much to ask that if we go to the trouble of collecting an infinite amount of data that we get the right answer
- The LLN basically states that the sample mean is consistent
- The sample variance and the sample standard deviation are consistent as well
- Recall also that the sample mean and the sample variance are unbiased as well
- (The sample standard deviation is biased, by the way)

# The Central Limit Theorem


- The **Central Limit Theorem** (CLT) is one of the most important theorems in statistics
- For our purposes, the CLT states that the distribution of averages of iid variables, properly normalized, becomes that of a standard normal as the sample size increases
- The CLT applies in an endless variety of settings
- Let  $X_1, \dots, X_n$  be a collection of iid random variables with mean  $\mu$  and variance  $\sigma^2$
- Let  $\bar{X}_n$  be their sample average
- Then  $\frac{\bar{X}_n - \mu}{\sigma / \sqrt{n}}$  has a distribution like that of a standard normal for large  $n$ .
- Remember the form  $\frac{\bar{X}_n - \mu}{\sigma / \sqrt{n}} = \frac{\text{Estimate} - \text{Mean of estimate}}{\text{Std. Err. of estimate}}$ .
- Usually, replacing the standard error by its estimated value doesn't change the CLT



# Example

- Simulate a standard normal random variable by rolling  $n$  (six sided)
- Let  $X_i$  be the outcome for die  $i$
- Then note that  $\mu = E[X_i] = 3.5$
- $\text{Var}(X_i) = 2.92$
- $\text{SE } \sqrt{2.92 / n} = 1.71 / \sqrt{n}$
- Standardized mean  $\frac{\bar{X}_n - 3.5}{1.71/\sqrt{n}}$


# Simulation of mean of $n$ dice




# Coin CLT

- Let  $X_i$  be the 0 or 1 result of the  $i^{\text{th}}$  flip of a possibly unfair coin
  - The sample proportion, say  $\hat{p}$ , is the average of the coin flips
  - $E[X_i] = p$  and  $\text{Var}(X_i) = p(1-p)$
  - Standard error of the mean is  $\sqrt{p(1-p)/n}$
  - Then  $\frac{\hat{p} - p}{\sqrt{p(1-p)/n}}$  will be approximately normally distributed


sample mean of 1 obs




sample mean of 10 obs




sample mean of 20 obs




sample mean of 1 obs



sample mean of 10 obs



sample mean of 20 obs



# CLT in practice

- In practice the CLT is mostly useful as an approximation  $P\left(\frac{\bar{X}_n - \mu}{\sigma / \sqrt{n}} \leq z\right) \approx \Phi(z)$ .
- Recall 1.96 is a good approximation to the .975<sup>th</sup> quantile of the standard normal
- Consider 
$$.95 \approx P\left(-1.96 \leq \frac{\bar{X}_n - \mu}{\sigma / \sqrt{n}} \leq 1.96\right) \iff .95 = P\left(\bar{X}_n + 1.96 \frac{\sigma}{\sqrt{n}} \geq \mu \geq \bar{X}_n - 1.96 \frac{\sigma}{\sqrt{n}}\right)$$

# Confidence intervals

- Therefore, according to the CLT, the probability that the random interval  $\bar{X}_n \pm z_{1-\alpha/2} \sigma / \sqrt{n}$  contains  $\mu$  is approximately  $100(1-\alpha)\%$ , where  $z_{1-\alpha/2}$  is the  $1-\alpha/2$  quantile of the standard normal distribution
- This is called a  $100(1 - \alpha)\%$  **confidence interval** for  $\mu$
- We can replace the unknown  $\sigma$  with  $s$

# Give a confidence interval for the average height of sons

in Galton's data

```
library(UsingR)
data(father.son)
x <- father.son$height
(mean(x) + c(-1, 1) * qnorm(0.975) * sd(x)/sqrt(length(x)))/12
```

```
## [1] 5.710 5.738
```

# Sample proportions

- In the event that each  $X_i$  is 0 or 1 with common success probability  $p$  then  $\sigma^2 = p(1 - p)$
- The interval takes the form  $\hat{p} \pm z_{1 - \alpha/2} \sqrt{\frac{p(1 - p)}{n}}$
- Replacing  $p$  by  $\hat{p}$  in the standard error results in what is called a Wald confidence interval for  $p$
- Also note that  $p(1-p) \leq 1/4$  for  $0 \leq p \leq 1$
- Let  $\alpha = .05$  so that  $z_{1 - \alpha/2} = 1.96 \approx 2$  then  $2 \sqrt{\frac{p(1 - p)}{n}} \leq 2 \sqrt{\frac{1}{4n}} = \frac{1}{\sqrt{n}}$
- Therefore  $\hat{p} \pm \frac{1}{\sqrt{n}}$  is a quick CI estimate for  $p$



# Example

- Your campaign advisor told you that in a random sample of 100 likely voters, 56 intent to vote for you.
  - Can you relax? Do you have this race in the bag?
  - Without access to a computer or calculator, how precise is this estimate?
- $1/\sqrt{100} = .1$  so a back of the envelope calculation gives an approximate 95% interval of  $(0.46, 0.66)$ 
  - Not enough for you to relax, better go do more campaigning!
- Rough guidelines, 100 for 1 decimal place, 10,000 for 2, 1,000,000 for 3.

```
round(1/sqrt(10^(1:6)), 3)
```

```
## [1] 0.316 0.100 0.032 0.010 0.003 0.001
```

# Poisson interval

- A nuclear pump failed 5 times out of 94.32 days, give a 95% confidence interval for the failure rate per day?
- $X \sim \text{Poisson}(\lambda t)$ .
- Estimate  $\hat{\lambda} = X/t$
- $\text{Var}(\hat{\lambda}) = \lambda / t$   $\frac{\hat{\lambda} - \lambda}{\sqrt{\hat{\lambda} / t}} = \frac{X - t\lambda}{\sqrt{X}} \rightarrow N(0,1)$
- This isn't the best interval.
  - There are better asymptotic intervals.
  - You can get an exact CI in this case.

## R code

```
x <- 5
t <- 94.32
lambda <- x/t
round(lambda + c(-1, 1) * qnorm(0.975) * sqrt(lambda/t), 3)
```

```
## [1] 0.007 0.099
```

```
poisson.test(x, T = 94.32)$conf
```

```
## [1] 0.01721 0.12371
## attr("conf.level")
## [1] 0.95
```

# In the regression class

```
exp(confint(glm(x ~ 1 + offset(log(t)), family = poisson(link = log))))
```

```
## Waiting for profiling to be done...
```

```
## 2.5 % 97.5 %
```

```
## 0.01901 0.11393
```